

Pipeline output 2.51 B tokens

**Source
narratives**
104 M
deduplicated

*19.1× more redundant
than source*

Copied context – redundant (1,695 M)

**79.4%
redundant**

Duplicated generated – redundant (298 M)

Scaffold (245 M)

**9.7%
scaffold**

Unique generated – trainable (168 M)

Unique source – trainable (104 M)

**10.9%
trainable**

Redundant content is 19.1× the unique source (1.99 B vs 104 M)

From 104 M source tokens to 2.51 B output tokens: only 10.9% is trainable-unique content

